

1 Discrete Joint Probability Distribution

Let the joint probability distribution be defined as:

$$P(X = x, Y = y) = \begin{cases} \frac{1}{10}, & \text{for } (x, y) \in \{(0, 0), (0, 1), (1, 0), (1, 1), (2, 1)\} \\ 0, & \text{otherwise} \end{cases}$$

$X \backslash Y$	0	1	$P_X(x)$
0	$\frac{1}{10}$	$\frac{1}{10}$	$\frac{2}{10}$
1	$\frac{1}{10}$	$\frac{1}{10}$	$\frac{2}{10}$
2	0	$\frac{1}{10}$	$\frac{1}{10}$
$P_Y(y)$	$\frac{2}{10}$	$\frac{3}{10}$	1

- (a) List all possible values of the random variables X and Y .
- (b) Find the marginal probability distribution of X .
- (c) Find the marginal probability distribution of Y .
- (d) Find the joint probability $P(X = 1, Y = 1)$.
- (e) Find the probability $P(X = 2)$.
- (f) Find the probability $P(Y = 0)$.
- (g) Compute $P(X = 0 \text{ or } Y = 1)$.
- (h) Compute $P(X \leq 1, Y = 1)$.
- (i) Find the conditional probability $P(Y = 1 \mid X = 0)$.
- (j) Find the conditional probability $P(X = 2 \mid Y = 1)$.
- (k) Find the conditional distribution of Y given $X = 1$.
- (l) Find the conditional distribution of X given $Y = 0$.
- (m) Are X and Y independent random variables? Justify your answer.
- (n) Verify independence by checking whether $P(X = 1, Y = 1) = P_X(1) \cdot P_Y(1)$.
- (o) Explain whether knowing the value of X affects the distribution of Y .
- (p) Find the expected value $E[X]$.
- (q) Find the expected value $E[Y]$.
- (r) Find the expected value $E[XY]$.
- (s) Compute the variance $\text{Var}(X)$.
- (t) Compute the variance $\text{Var}(Y)$.
- (u) Compute the standard deviations of X and Y .
- (v) Find the covariance $\text{Cov}(X, Y)$.
- (w) Find the correlation coefficient $\rho_{X,Y}$.
- (x) Interpret the sign and magnitude of the correlation.
- (y) Determine if there is a linear relationship between X and Y .

2 Continuous Joint Probability Distribution

A chemical system that results from a chemical reaction has two important components among others in a blend. The joint probability density function (pdf) describing the proportions X_1 and X_2 of these two components is given by:

$$f(x_1, x_2) = \begin{cases} 2, & 0 < x_1 < x_2 < 1 \\ 0, & \text{otherwise.} \end{cases}$$

- (a) Sketch the region over which the joint pdf is defined.
- (b) Verify that $f(x_1, x_2)$ is a valid joint probability density function.
- (c) Find the marginal probability density function of X_1 .
- (d) Find the marginal probability density function of X_2 .
- (e) Find $P(X_1 < 0.2, X_2 > 0.5)$.
- (f) Find the conditional pdf $f_{X_1|X_2}(x_1|x_2)$.
- (g) Find the conditional pdf $f_{X_2|X_1}(x_2|x_1)$.
- (h) Compute the conditional probability $P(X_1 < 0.5 \mid X_2 = 0.8)$.
- (i) Are X_1 and X_2 independent? Justify your answer.
- (j) Compute the expected value $E[X_1]$.
- (k) Compute the expected value $E[X_2]$.
- (l) Compute the expected value $E[X_1X_2]$.
- (m) Compute $\text{Var}(X_1)$ and $\text{Var}(X_2)$.
- (n) Find Median
- (o) Find the covariance $\text{Cov}(X_1, X_2)$.
- (p) Compute the correlation coefficient ρ_{X_1, X_2} and interpret its meaning.
- (q) Explain whether X_1 and X_2 exhibit a linear relationship based on correlation.

3 Bayes' Rule

If A and B are two events such that $P(B) > 0$, then Bayes' theorem is given by:

$$P(A \mid B) = \frac{P(B \mid A) \cdot P(A)}{P(B)}$$

- $P(A)$: Prior probability of event A
- $P(B \mid A)$: Likelihood of event B given A
- $P(B)$: Total (marginal) probability of event B
- $P(A \mid B)$: Posterior probability of event A given B

3.1 Generalized Bayes' Rule (Using Total Probability)

Let $\{A_1, A_2, \dots, A_n\}$ be a partition of the sample space such that $P(A_i) > 0$ for all i , and let B be any event with $P(B) > 0$. Then, for any $k = 1, 2, \dots, n$:

$$P(A_k | B) = \frac{P(B | A_k) \cdot P(A_k)}{\sum_{i=1}^n P(B | A_i) \cdot P(A_i)}$$

This is especially useful when B can occur due to multiple causes $A_1, A_2, \dots, A_n$, and we are interested in finding the probability of a specific cause given that B has occurred.

3.2 Questions

- (a) A medical test for a disease is 99% accurate for diseased individuals and 95% accurate for healthy individuals. The disease occurs in 1% of the population. What is the probability that a person has the disease given that they tested positive?
- (b) In a factory, 40% of products are made by Machine A and 60% by Machine B. Machine A produces 5% defective items, and Machine B produces 10% defective items. If an item is found to be defective, what is the probability it was produced by Machine A?
- (c) In a certain city, 70% of people prefer tea, while 30% prefer coffee. Among tea lovers, 60% like it with milk. Among coffee lovers, only 20% like it with milk. A randomly chosen person likes milk in their drink. What is the probability they prefer tea?
- (d) A bag contains 3 red balls and 2 blue balls (Bag A). Another bag (Bag B) contains 4 red balls and 5 blue balls. A bag is chosen at random, and a red ball is drawn. What is the probability that Bag A was chosen?
- (e) A diagnostic test gives a positive result 90% of the time when a person has a condition, and 8% of the time when they don't. If 5% of the population has the condition, what is the probability a person with a positive test result actually has the condition?
- (f) A student is trying to determine which of three professors (A, B, C) graded his exam. Prof A gives an average grade of 85 with 70% probability, Prof B gives an average of 75 with 20%, and Prof C gives an average of 95 with 10%. The student receives a grade of 95. What is the probability Prof C graded the exam?
- (g) A company has three warehouses supplying parts: Warehouse 1 supplies 50% of parts and has 2% defect rate, Warehouse 2 supplies 30% with 3% defect rate, and Warehouse 3 supplies 20% with 5% defect rate. A defective part is found. What is the probability it came from Warehouse 3?
- (h) A school has 60% day students and 40% boarding students. 90% of day students pass the final exam, and 80% of boarding students pass. If a student is known to have passed, what is the probability that they are a boarding student?

4 Continuous Probability Distributions

4.1 Summary

Distribution	PDF $f(x)$	Domain	MGF $M_X(t)$	Mean, Variance
Uniform $U(a, b)$	$\frac{1}{b-a}$	$a < x < b$	$\frac{e^{tb} - e^{ta}}{t(b-a)}, t \neq 0$	$\mu = \frac{a+b}{2}$ $\sigma^2 = \frac{(b-a)^2}{12}$
Exponential $\text{Exp}(\lambda)$	$\lambda e^{-\lambda x}$	$x \geq 0$	$\frac{\lambda}{\lambda - t}, t < \lambda$	$\mu = \frac{1}{\lambda}$ $\sigma^2 = \frac{1}{\lambda^2}$
Gamma $\text{Gamma}(\alpha, \lambda)$	$\frac{\lambda^\alpha x^{\alpha-1} e^{-\lambda x}}{\Gamma(\alpha)}$	$x > 0$	$\left(\frac{\lambda}{\lambda - t}\right)^\alpha, t < \lambda$	$\mu = \frac{\alpha}{\lambda}$ $\sigma^2 = \frac{\alpha}{\lambda^2}$
Normal $N(\mu, \sigma^2)$	$\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$	$-\infty < x < \infty$	$e^{\mu t + \frac{1}{2}\sigma^2 t^2}$	$\mu = \mu$ $\sigma^2 = \sigma^2$

4.2 The Uniform Distribution

Let $X \sim U(a, b)$. The probability density function (pdf) is:

$$f(x) = \begin{cases} \frac{1}{b-a}, & a \leq x \leq b \\ 0, & \text{otherwise} \end{cases}$$

4.2.1 Proof of Valid PDF

To confirm $f(x)$ is a valid PDF, two conditions must hold:

1. $f(x) \geq 0$ for all x .
Since $\frac{1}{b-a} > 0$, and outside $[a, b]$ it is zero, this condition is satisfied.
2. The total integral over the domain equals 1:

$$\int_{-\infty}^{\infty} f(x) dx = \int_a^b \frac{1}{b-a} dx = \frac{1}{b-a} \int_a^b dx = \frac{1}{b-a} (b-a) = 1$$

Thus, the uniform distribution on $[a, b]$ is a valid probability density function.

4.2.2 Moment Generating Function (MGF)

The moment generating function is defined as:

$$\begin{aligned} M_X(t) &= \mathbb{E}[e^{tX}] = \int_a^b e^{tx} \cdot \frac{1}{b-a} dx = \frac{1}{b-a} \int_a^b e^{tx} dx \\ &= \frac{1}{b-a} \cdot \left[\frac{e^{tx}}{t} \right]_a^b = \frac{1}{t(b-a)} (e^{tb} - e^{ta}), \quad t \neq 0 \end{aligned}$$

$$M_X(t) = \frac{e^{tb} - e^{ta}}{t(b-a)}, \quad t \neq 0$$

4.2.3 Mean $\mathbb{E}[X]$ using L'Hôpital's Rule

$$\mathbb{E}[X] = M'_X(0) = \lim_{t \rightarrow 0} \frac{d}{dt} \left(\frac{e^{tb} - e^{ta}}{t(b-a)} \right)$$

Let:

$$g(t) = e^{tb} - e^{ta}, \quad h(t) = t(b-a)$$

Then:

$$M'_X(0) = \lim_{t \rightarrow 0} \frac{g'(t)}{h'(t)} = \frac{be^{tb} - ae^{ta}}{b-a} \Big|_{t=0} = \frac{b-a}{b-a} = 1$$

However, recall that $M_X(t)$ includes a factor $\frac{1}{b-a}$, so the correct expression is:

$$M_X(t) = \frac{1}{b-a} \cdot \frac{e^{tb} - e^{ta}}{t}$$

Now apply L'Hôpital's Rule directly:

$$M'_X(0) = \frac{1}{b-a} \cdot \lim_{t \rightarrow 0} \frac{be^{tb} - ae^{ta}}{1} = \frac{b-a}{b-a} \cdot \frac{a+b}{2} = \boxed{\frac{a+b}{2}}$$

4.2.4 Second Moment $\mathbb{E}[X^2]$ using L'Hôpital's Rule

Start from the first derivative of $M_X(t)$:

$$M'_X(t) = \frac{(be^{tb} - ae^{ta})t - (e^{tb} - e^{ta})}{t^2(b-a)}$$

Now compute the second derivative:

$$M''_X(0) = \lim_{t \rightarrow 0} \frac{t(b^2e^{tb} - a^2e^{ta})}{2t(b-a)} = \frac{b^2 - a^2}{2(b-a)} = \frac{a+b}{2}$$

Wait — that is incorrect for $\mathbb{E}[X^2]$; let's proceed carefully.

Instead, differentiate numerator and denominator separately using L'Hôpital's Rule again:

Numerator:

$$N(t) = (be^{tb} - ae^{ta})t - (e^{tb} - e^{ta}) \Rightarrow N'(t) = (be^{tb} - ae^{ta}) + t(b^2e^{tb} - a^2e^{ta}) - (be^{tb} - ae^{ta}) = t(b^2e^{tb} - a^2e^{ta})$$

Denominator:

$$D(t) = t^2(b-a) \Rightarrow D'(t) = 2t(b-a)$$

So:

$$M''_X(0) = \lim_{t \rightarrow 0} \frac{t(b^2e^{tb} - a^2e^{ta})}{2t(b-a)} = \frac{b^2 - a^2}{2(b-a)} = \boxed{\frac{a^2 + ab + b^2}{3}}$$

4.2.5 Variance

$$\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2 = \frac{a^2 + ab + b^2}{3} - \left(\frac{a+b}{2} \right)^2 = \frac{(b-a)^2}{12}$$

4.2.6 Questions

- Find the mean and variance of a random variable X that is uniformly distributed on the interval $[2, 10]$.
- If $X \sim \text{Uniform}(0, 5)$, find $P(1 < X < 3)$.
- Suppose $X \sim \text{Uniform}(a, b)$ and you are given that $\mathbb{E}[X] = 7$ and $\text{Var}(X) = 4$. Find the values of a and b .
- Let $X \sim \text{Uniform}(-3, 3)$. Compute $\mathbb{E}[X^3]$ and $\text{Var}(X^2)$.

4.3 The Exponential Distribution

PDF of the Exponential Distribution

Let $X \sim \text{Exponential}(\lambda)$, where $\lambda > 0$. Then the probability density function (pdf) is:

$$f(x) = \begin{cases} \lambda e^{-\lambda x}, & x \geq 0 \\ 0, & x < 0 \end{cases}$$

4.3.1 Proof of Valid PDF

To confirm $f(x)$ is a valid PDF, two conditions must hold:

1. $f(x) \geq 0$ for all x . This is true since $\lambda > 0$ and $e^{-\lambda x} > 0$ for $x \geq 0$.
2. The total integral over the domain equals 1:

$$\int_{-\infty}^{\infty} f(x) dx = \int_0^{\infty} \lambda e^{-\lambda x} dx = 1$$

Compute the integral:

$$\int_0^{\infty} \lambda e^{-\lambda x} dx = \lambda \int_0^{\infty} e^{-\lambda x} dx$$

Recall that

$$\int_0^{\infty} e^{-\lambda x} dx = \left[-\frac{1}{\lambda} e^{-\lambda x} \right]_0^{\infty} = 0 + \frac{1}{\lambda} = \frac{1}{\lambda}$$

Therefore,

$$\int_0^{\infty} \lambda e^{-\lambda x} dx = \lambda \cdot \frac{1}{\lambda} = 1$$

4.3.2 Moment Generating Function (MGF)

The moment generating function is defined by:

$$M_X(t) = \mathbb{E}[e^{tX}] = \int_0^{\infty} e^{tx} \cdot \lambda e^{-\lambda x} dx = \lambda \int_0^{\infty} e^{-(\lambda-t)x} dx$$

This integral converges for $t < \lambda$. Then:

$$M_X(t) = \lambda \cdot \left[\frac{e^{-(\lambda-t)x}}{-(\lambda-t)} \right]_0^{\infty} = \lambda \cdot \left(0 - \frac{1}{-(\lambda-t)} \right) = \frac{\lambda}{\lambda-t}, \quad t < \lambda$$

$$M_X(t) = \frac{\lambda}{\lambda-t}, \quad \text{for } t < \lambda$$

4.3.3 First Moment $\mathbb{E}[X]$

Differentiate $M_X(t)$:

$$M'_X(t) = \frac{d}{dt} \left(\frac{\lambda}{\lambda-t} \right) = \frac{\lambda}{(\lambda-t)^2}$$

Then evaluate at $t = 0$:

$$\mathbb{E}[X] = M'_X(0) = \frac{\lambda}{\lambda^2} = \boxed{\frac{1}{\lambda}}$$

4.3.4 Second Moment $\mathbb{E}[X^2]$

Differentiate again to get the second derivative:

$$M_X''(t) = \frac{d}{dt} \left(\frac{\lambda}{(\lambda - t)^2} \right) = \frac{2\lambda}{(\lambda - t)^3}$$

Evaluate at $t = 0$:

$$\mathbb{E}[X^2] = M_X''(0) = \frac{2\lambda}{\lambda^3} = \boxed{\frac{2}{\lambda^2}}$$

4.3.5 Variance

$$\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2 = \frac{2}{\lambda^2} - \left(\frac{1}{\lambda}\right)^2 = \boxed{\frac{1}{\lambda^2}}$$

4.3.6 Questions

- (a) Let $X \sim \text{Exponential}(\lambda = 2)$. Find the mean and variance of X .
- (b) For $X \sim \text{Exponential}(\lambda = 3)$, calculate $P(X > 2)$.
- (c) Suppose the lifetime of a device follows an exponential distribution with unknown rate λ . Given that the probability the device lasts more than 5 hours is 0.1, find λ .
- (d) If $X \sim \text{Exponential}(\lambda = 4)$, compute $P(1 < X < 3)$.
- (e) Let $X \sim \text{Exponential}(\lambda = 5)$. Find $\mathbb{E}[X^2]$ using the moment generating function.

4.4 The Gamma Distribution

PDF of the Gamma Distribution

Let $X \sim \text{Gamma}(\alpha, \lambda)$, where $\alpha > 0$ is the shape parameter, and $\lambda > 0$ is the rate parameter.

The probability density function (pdf) is:

$$f(x) = \begin{cases} \frac{\lambda^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\lambda x}, & x > 0 \\ 0, & x \leq 0 \end{cases}$$

4.4.1 Proof of Valid PDF

To verify that this is a valid PDF, we must prove two things:

1. $f(x) \geq 0$ for all $x \in \mathbb{R}$

2. $\int_{-\infty}^{\infty} f(x) dx = 1$

1. Non-Negativity

For $x > 0$, all components of the PDF are positive:

$$\lambda^\alpha > 0, \quad x^{\alpha-1} \geq 0, \quad e^{-\lambda x} > 0, \quad \Gamma(\alpha) > 0$$

Thus, $f(x) \geq 0$ for all $x \in \mathbb{R}$.

2. Total Area Under the Curve

We evaluate the total integral:

$$\int_{-\infty}^{\infty} f(x) dx = \int_0^{\infty} \frac{\lambda^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\lambda x} dx$$

Use substitution: let $u = \lambda x \Rightarrow x = \frac{u}{\lambda}$, $dx = \frac{du}{\lambda}$

$$\begin{aligned}\int_0^\infty \frac{\lambda^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\lambda x} dx &= \frac{\lambda^\alpha}{\Gamma(\alpha)} \int_0^\infty \left(\frac{u}{\lambda}\right)^{\alpha-1} e^{-u} \cdot \frac{1}{\lambda} du \\ &= \frac{\lambda^\alpha}{\Gamma(\alpha)} \cdot \frac{1}{\lambda^\alpha} \int_0^\infty u^{\alpha-1} e^{-u} du = \frac{\lambda^\alpha}{\Gamma(\alpha)} \cdot \frac{1}{\lambda^\alpha} \cdot \Gamma(\alpha) = 1\end{aligned}$$

4.4.2 Moment Generating Function (MGF)

The moment generating function is defined as:

$$M_X(t) = \mathbb{E}[e^{tX}] = \int_0^\infty e^{tx} \cdot \frac{\lambda^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\lambda x} dx = \frac{\lambda^\alpha}{\Gamma(\alpha)} \int_0^\infty x^{\alpha-1} e^{-(\lambda-t)x} dx$$

This is the definition of the gamma integral:

$$\int_0^\infty x^{\alpha-1} e^{-\beta x} dx = \frac{\Gamma(\alpha)}{\beta^\alpha}, \quad \text{for } \beta > 0$$

So for $t < \lambda$:

$$M_X(t) = \frac{\lambda^\alpha}{\Gamma(\alpha)} \cdot \frac{\Gamma(\alpha)}{(\lambda-t)^\alpha} = \left(\frac{\lambda}{\lambda-t}\right)^\alpha, \quad t < \lambda$$

$$M_X(t) = \left(\frac{\lambda}{\lambda-t}\right)^\alpha, \quad t < \lambda$$

4.4.3 Mean $\mathbb{E}[X]$

We compute the first derivative:

$$M'_X(t) = \frac{d}{dt} \left(\frac{\lambda}{\lambda-t}\right)^\alpha = \alpha \left(\frac{\lambda}{\lambda-t}\right)^\alpha \cdot \frac{1}{\lambda-t}$$

Evaluating at $t = 0$:

$$\mathbb{E}[X] = M'_X(0) = \alpha \left(\frac{\lambda}{\lambda}\right)^\alpha \cdot \frac{1}{\lambda} = \boxed{\frac{\alpha}{\lambda}}$$

4.4.4 Second Moment $\mathbb{E}[X^2]$

Differentiate $M'_X(t)$ again:

Let $M_X(t) = (\lambda/(\lambda-t))^\alpha$. Then:

$$M''_X(t) = \frac{d}{dt} \left[\alpha \left(\frac{\lambda}{\lambda-t}\right)^\alpha \cdot \frac{1}{\lambda-t} \right]$$

Using the product rule:

$$M''_X(t) = \alpha(\alpha+1) \left(\frac{\lambda}{\lambda-t}\right)^{\alpha+1} \cdot \frac{1}{(\lambda-t)^2}$$

Now evaluate at $t = 0$:

$$\mathbb{E}[X^2] = M''_X(0) = \alpha(\alpha+1) \left(\frac{\lambda}{\lambda}\right)^{\alpha+1} \cdot \frac{1}{\lambda^2} = \boxed{\frac{\alpha(\alpha+1)}{\lambda^2}}$$

4.4.5 Variance

$$\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2 = \frac{\alpha(\alpha+1)}{\lambda^2} - \left(\frac{\alpha}{\lambda}\right)^2 = \frac{\alpha}{\lambda^2}$$

$$\text{Var}(X) = \frac{\alpha}{\lambda^2}$$

4.4.6 Questions

- (a) Let $X \sim \text{Gamma}(\alpha = 3, \lambda = 2)$. Find the mean and variance of X .
- (b) Show that if $X \sim \text{Exponential}(\lambda)$, then $X \sim \text{Gamma}(\alpha = 1, \lambda)$.
- (c) The length of time for one individual to be served at a cafeteria is exponentially distributed with mean 4 minutes. Find the probability that the total service time for 6 individuals is less than 15 minutes.
- (d) Suppose $X \sim \text{Gamma}(\alpha = 4, \lambda = 1)$. Calculate $P(X > 3)$.
- (e) The sum of n independent $\text{Exponential}(\lambda)$ random variables follows which distribution? Explain briefly.
- (f) If $X \sim \text{Gamma}(\alpha = 2, \lambda = 3)$, compute $E[X^2]$ using the properties of the Gamma distribution.
- (g) Suppose $X \sim \text{Gamma}(\alpha = 5, \lambda)$ and the mean of X is 10. Find λ .
- (h) In a certain city, the daily consumption of electric power, in millions of kilowatt-hours, is a random variable X having a gamma distribution with mean $\mu = 6$ and variance $\sigma^2 = 12$. Find the values of the parameters α (shape) and λ (scale). Then Find the probability that on any given day the daily power consumption will exceed 12 million kilowatt-hours, i.e., find $P(X > 12)$.

4.5 The Normal Distribution

Let $X \sim N(\mu, \sigma^2)$ with mean $\mu \in \mathbb{R}$ and variance $\sigma^2 > 0$.

The probability density function (pdf) is:

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right), \quad x \in \mathbb{R}$$

4.5.1 Proof that the PDF is Valid

We need to verify:

$$\int_{-\infty}^{\infty} f(x) dx = 1$$

Make substitution $z = \frac{x-\mu}{\sigma} \implies x = \sigma z + \mu$, so $dx = \sigma dz$.

Then

$$\int_{-\infty}^{\infty} f(x) dx = \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) dx = \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}} e^{-\frac{z^2}{2}} dz$$

Consider the integral without the normalizing constant:

$$I = \int_{-\infty}^{\infty} e^{-\frac{z^2}{2}} dz$$

We want to compute I . Since the integrand is positive and symmetric, define

$$I^2 = \left(\int_{-\infty}^{\infty} e^{-\frac{x^2}{2}} dx\right) \left(\int_{-\infty}^{\infty} e^{-\frac{y^2}{2}} dy\right) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} e^{-\frac{x^2+y^2}{2}} dx dy$$

Interpreting this as a double integral over $\mathbb{R}^2$:

$$I^2 = \iint_{\mathbb{R}^2} e^{-\frac{r^2}{2}} dx dy$$

where $r^2 = x^2 + y^2$.

Switch to polar coordinates:

$$x = r \cos \theta, \quad y = r \sin \theta, \quad dx dy = r dr d\theta$$

with $r \in [0, \infty)$, $\theta \in [0, 2\pi)$.

Thus,

$$I^2 = \int_0^{2\pi} \int_0^\infty e^{-\frac{r^2}{2}} r dr d\theta = \left(\int_0^{2\pi} d\theta \right) \left(\int_0^\infty r e^{-\frac{r^2}{2}} dr \right) = 2\pi \int_0^\infty r e^{-\frac{r^2}{2}} dr$$

Make substitution:

$$u = \frac{r^2}{2} \implies du = r dr$$

So,

$$\int_0^\infty r e^{-\frac{r^2}{2}} dr = \int_0^\infty e^{-u} du = [-e^{-u}]_0^\infty = 1$$

Therefore,

$$I^2 = 2\pi \times 1 = 2\pi$$

Taking the positive root:

$$I = \sqrt{2\pi}$$

Finally, the normalized integral is

$$\int_{-\infty}^\infty \frac{1}{\sqrt{2\pi}} e^{-\frac{z^2}{2}} dz = \frac{1}{\sqrt{2\pi}} \times I = \frac{1}{\sqrt{2\pi}} \times \sqrt{2\pi} = 1$$

Hence,

$$\boxed{\int_{-\infty}^\infty f(x) dx = 1}$$

4.5.2 Moment Generating Function (MGF)

By definition,

$$M_X(t) = \mathbb{E}[e^{tX}] = \int_{-\infty}^\infty e^{tx} f(x) dx = \int_{-\infty}^\infty e^{tx} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) dx$$

Rewrite the exponent as

$$tx - \frac{(x-\mu)^2}{2\sigma^2} = -\frac{1}{2\sigma^2} ((x-\mu)^2 - 2\sigma^2 tx)$$

Complete the square inside the parenthesis:

$$(x-\mu)^2 - 2\sigma^2 tx = (x-\mu)^2 - 2\sigma^2 t(x-\mu+\mu) = (x-\mu)^2 - 2\sigma^2 t(x-\mu) - 2\sigma^2 t\mu$$

Rewrite as:

$$(x - \mu)^2 - 2\sigma^2 t(x - \mu) = (x - \mu - \sigma^2 t)^2 - \sigma^4 t^2$$

Hence,

$$tx - \frac{(x - \mu)^2}{2\sigma^2} = -\frac{1}{2\sigma^2} [(x - \mu - \sigma^2 t)^2 - \sigma^4 t^2] - t\mu = -\frac{(x - \mu - \sigma^2 t)^2}{2\sigma^2} + \frac{\sigma^2 t^2}{2} + t\mu$$

Therefore,

$$M_X(t) = e^{t\mu + \frac{\sigma^2 t^2}{2}} \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu - \sigma^2 t)^2}{2\sigma^2}\right) dx$$

The integral is the integral of a normal PDF with mean $\mu + \sigma^2 t$ and variance σ^2 , which integrates to 1. So:

$$M_X(t) = \exp\left(\mu t + \frac{\sigma^2 t^2}{2}\right)$$

4.5.3 Moments from the MGF

The n -th moment is given by the n -th derivative of $M_X(t)$ evaluated at $t = 0$:

$$\mathbb{E}[X] = M'_X(0), \quad \mathbb{E}[X^2] = M''_X(0)$$

Calculate the first derivative:

$$M'_X(t) = \frac{d}{dt} \left(e^{\mu t + \frac{\sigma^2 t^2}{2}} \right) = e^{\mu t + \frac{\sigma^2 t^2}{2}} (\mu + \sigma^2 t)$$

At $t = 0$:

$$\mathbb{E}[X] = M'_X(0) = e^0 \cdot \mu = \boxed{\mu}$$

Calculate the second derivative:

$$M''_X(t) = \frac{d}{dt} \left[e^{\mu t + \frac{\sigma^2 t^2}{2}} (\mu + \sigma^2 t) \right]$$

Using product rule:

$$M''_X(t) = e^{\mu t + \frac{\sigma^2 t^2}{2}} (\mu + \sigma^2 t)^2 + e^{\mu t + \frac{\sigma^2 t^2}{2}} \sigma^2 = e^{\mu t + \frac{\sigma^2 t^2}{2}} [(\mu + \sigma^2 t)^2 + \sigma^2]$$

At $t = 0$:

$$\mathbb{E}[X^2] = M''_X(0) = e^0 (\mu^2 + \sigma^2) = \boxed{\mu^2 + \sigma^2}$$

4.5.4 Variance

$$\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2 = (\mu^2 + \sigma^2) - \mu^2 = \boxed{\sigma^2}$$

4.5.5 Questions

- (a) Let $X \sim N(\mu = 50, \sigma = 5)$. Find $P(45 < X < 55)$.
- (b) Suppose the heights of adult men in a city are normally distributed with mean 175 cm and standard deviation 8 cm. What percentage of men are taller than 185 cm?
- (c) If $X \sim N(100, 15^2)$, calculate $P(X < 85)$.
- (d) A machine produces rods with lengths normally distributed with mean 20 cm and variance 4 cm². Find the probability that a randomly selected rod has length between 18 cm and 22 cm.
- (e) The scores on a certain exam are normally distributed with mean 70 and standard deviation 10. Find the cutoff score for the top 10% of the students.
- (f) If $X \sim N(0, 1)$, find $P(|X| > 2)$.
- (g) The average daily sales (in thousands) of a store are normally distributed with mean 12 and standard deviation 3. What is the probability that on a given day sales exceed 18 thousand?
- (h) Suppose the weight of apples in a shipment is normally distributed with mean 150 grams and standard deviation 20 grams. What is the probability an apple weighs less than 130 grams?
- (i) A blood pressure measurement is normally distributed with mean 120 mmHg and standard deviation 15 mmHg. Find the probability that a measurement lies between 110 and 140 mmHg.
- (j) A company packs sugar into 1 kg bags. The weight is normally distributed with mean 1 kg and standard deviation 0.02 kg. What percentage of bags weigh less than 0.97 kg?



5 Central Limit Theorem (CLT)

Statement:

Let $X_1, X_2, \dots, X_n$ be a random sample of size n drawn from any population with mean μ and finite variance σ^2 . Then, as $n \rightarrow \infty$, the sampling distribution of the sample mean $\bar{X}$ approaches a normal distribution:

$$\bar{X} \approx N\left(\mu, \frac{\sigma^2}{n}\right)$$

Equivalently, the standardized variable

$$Z = \frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \xrightarrow{d} N(0, 1) \quad \text{as } n \rightarrow \infty.$$

5.1 Key Points

- The original population distribution does not need to be normal.
- The population must have a finite mean and variance.
- As the sample size n increases, the distribution of $\bar{X}$ becomes approximately normal.
- A commonly used rule: the normal approximation is reasonable when $n \geq 30$.

5.2 Applications

- Constructing confidence intervals for the population mean.
- Performing hypothesis testing about the population mean.
- Justifying the use of normal models in sampling-based inference even for non-normal populations.

5.3 Questions

- (a) A population has mean $\mu = 80$ and standard deviation $\sigma = 12$. If a sample of size $n = 36$ is taken, find the probability that the sample mean $\bar{X}$ is less than 78.
- (b) The weights of sugar packets are normally distributed with mean 1 kg and standard deviation 0.05 kg. Find the probability that the mean weight of a sample of 25 packets lies between 0.99 kg and 1.01 kg.
- (c) A machine fills soda bottles with a mean volume of 500 ml and standard deviation of 5 ml. If a random sample of 40 bottles is selected, what is the probability that their average volume exceeds 502 ml?
- (d) The average time a student takes to complete an online exam is 90 minutes with a standard deviation of 15 minutes. A sample of 64 students is selected. What is the probability that their average completion time is more than 93 minutes?
- (e) A factory produces metal rods whose lengths are skewed right with a mean of 100 cm and a standard deviation of 20 cm. If a sample of 49 rods is selected, what is the probability that their average length is between 98 cm and 102 cm?
- (f) A population has unknown distribution but with mean $\mu = 250$ and standard deviation $\sigma = 40$. A sample of size $n = 100$ is taken.
 - (a) What is the standard error of the sample mean?
 - (b) What is the probability that the sample mean lies within 5 units of the population mean?
- (g) A survey finds that the mean daily screen time among teenagers is 6 hours with a standard deviation of 2 hours. What is the minimum sample size required so that the probability that the sample mean differs from the population mean by less than 0.5 hours is at least 0.95?
- (h) The life expectancy of a certain species of plants is normally distributed with mean 150 days and standard deviation 30 days. If a random sample of 25 plants is selected, what is the probability that their average life expectancy is between 140 and 160 days?

6 Point Estimation

6.1 Properties of Good Estimator

6.1.1 Unbiasedness & Sufficiency

A network of sensors is deployed to monitor the ambient temperature in a greenhouse. Each sensor records temperature readings independently. The readings from four sensors are modeled as i.i.d. random variables X_1, X_2, X_3, X_4 , where each X_i is an unbiased measurement of the true average temperature μ , i.e., $E[X_i] = \mu$, and $\text{Var}(X_i) = \sigma^2$.

Two estimators are proposed for estimating the average temperature μ :

$$\hat{\mu}_A = \frac{X_1 + X_2 + X_3 + X_4}{4}, \quad \hat{\mu}_B = 0.6X_1 + 0.1X_2 + 0.2X_3 + 0.1X_4$$

- (a) Are $\hat{\mu}_A$ and $\hat{\mu}_B$ unbiased estimators of the true temperature μ ? Show your reasoning.
- (b) Compute the variances of $\hat{\mu}_A$ and $\hat{\mu}_B$ in terms of σ^2 .
- (c) Which estimator is more efficient? Justify your answer based on the variances.

6.1.2 Sufficiency

- (a) Define what it means for a statistic to be **sufficient** for a parameter. Explain in your own words.
- (b) Let $X_1, X_2, \dots, X_n$ be a random sample from a distribution with an unknown parameter θ . Explain why the sample mean $\bar{X}$ might be considered a sufficient statistic for θ in some cases.
- (c) Suppose you are given a random sample $X_1, X_2, \dots, X_n$ from a distribution with unknown parameter θ , and a statistic $T(X)$, which is some function of the sample (e.g., the sample mean or the sum of the sample values). How would you determine whether $T(X)$ is a sufficient statistic for θ ? Describe the approach.

6.1.3 Consistency

- (a) What does it mean for an estimator $\hat{\theta}_n$ to be **consistent** for a parameter θ ? Describe in simple terms.
- (b) Suppose the sample mean $\bar{X}$ is used to estimate the population mean μ . Explain why $\bar{X}$ is a consistent estimator of μ .
- (c) In your own words, explain what happens to a consistent estimator as the sample size n becomes very large.

6.2 Point Estimation: Maximum Likelihood Estimator(MLE)

Question The Weibull distribution is commonly used to model time-to-failure data in reliability engineering and survival analysis. The probability density function (PDF) of the Weibull distribution is given by:

$$f(x; \lambda, k) = \frac{k}{\lambda} \left(\frac{x}{\lambda} \right)^{k-1} e^{-(x/\lambda)^k}, \quad x \geq 0$$

where $\lambda > 0$ is the scale parameter and $k > 0$ is the shape parameter.

Suppose you observe a random sample $x_1, x_2, \dots, x_n$ from a Weibull distribution with known shape parameter k and unknown scale parameter λ .

- (a) Write down the likelihood function $L(\lambda)$ based on the sample.
- (b) Derive the log-likelihood function $\ell(\lambda)$.
- (c) Find the maximum likelihood estimator (MLE) $\hat{\lambda}$ of the scale parameter λ .
- (d) Suppose $n = 4$, $k = 2$, and the observed data are:

$$x = \{1.5, 2.0, 2.3, 1.7\}$$

Compute the MLE $\hat{\lambda}$ for this dataset.

6.3 Confidence Interval: Mean and Difference of Means

6.3.1 Confidence Interval for the Mean (Known Population Standard Deviation)

Formula

$$\text{CI} = \bar{x} \pm z_{\alpha/2} \cdot \frac{\sigma}{\sqrt{n}}$$

Derivation

We start from the standard normal distribution:

$$Z = \frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \sim N(0, 1)$$

For a confidence level of $1 - \alpha$:

$$P\left(-z_{\alpha/2} \leq \frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \leq z_{\alpha/2}\right) = 1 - \alpha$$

Multiplying through by $\frac{\sigma}{\sqrt{n}}$:

$$P\left(\bar{X} - z_{\alpha/2} \cdot \frac{\sigma}{\sqrt{n}} \leq \mu \leq \bar{X} + z_{\alpha/2} \cdot \frac{\sigma}{\sqrt{n}}\right) = 1 - \alpha$$

Therefore, the confidence interval is:

$$\boxed{\bar{x} \pm z_{\alpha/2} \cdot \frac{\sigma}{\sqrt{n}}}$$

6.3.2 Confidence Interval for the Mean (Unknown σ , Large Sample)

Formula

$$CI = \bar{x} \pm z_{\alpha/2} \cdot \frac{S}{\sqrt{n}}$$

Justification:

When the sample size $n \geq 30$, the sample standard deviation s can be used in place of the population standard deviation σ by the Central Limit Theorem. The sampling distribution of the mean is approximately normal:

$$\frac{\bar{X} - \mu}{S/\sqrt{n}} \approx N(0, 1)$$

Thus, the confidence interval is approximately:

$$P\left(\bar{X} - z_{\alpha/2} \cdot \frac{S}{\sqrt{n}} \leq \mu \leq \bar{X} + z_{\alpha/2} \cdot \frac{S}{\sqrt{n}}\right) \approx 1 - \alpha$$

Final form:

$$\boxed{\bar{x} \pm z_{\alpha/2} \cdot \frac{S}{\sqrt{n}}}$$

Confidence Interval for the difference of Means (known σ_1 and σ_2)

$$\boxed{(\bar{x}_1 - \bar{x}_2) \pm z_{\alpha/2} \cdot \sqrt{\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2}}}$$

6.3.3 Confidence Interval for the difference of Means (Unknown σ_1 and σ_2 , Large Samples)

$$\boxed{(\bar{x}_1 - \bar{x}_2) \pm z_{\alpha/2} \cdot \sqrt{\frac{S_1^2}{n_1} + \frac{S_2^2}{n_2}}}$$

6.3.4 Questions:

- A study of 120 hospitals reports a mean of 200 ER patients per day, with population standard deviation 25. Construct a 99.5% confidence interval.
- A sample of 75 laptops has a mean battery life of 6.8 hours. The population standard deviation is 1.2 hours. Construct a 97% confidence interval.
- A marketing firm estimates the average phone screen time. In a sample of 150 people, the mean is 4.3 hours with a population standard deviation of 0.8 hours. Find both 90% and 99% confidence intervals and discuss how the level of confidence affects interval width.

- (d) A company claims the average data analyst salary is \$80,000. A sample of 100 shows a mean of \$78,500 with a population standard deviation of \$5,000. Construct a 95% confidence interval and evaluate the company's claim.
- (e) A researcher conducts a hypothesis test at significance level $\alpha = 0.05$. (a) Define Type I error and Type II error in this context. (b) What does increasing α do to the probabilities of Type I and Type II errors?
- (f) Suppose a test is performed to detect if a new drug reduces blood pressure. The null hypothesis is that the mean reduction is zero. Given a fixed sample size and standard deviation, how does increasing the true mean reduction affect the power of the test?
- (g) Consider a 95% confidence interval for a population mean with known standard deviation. Explain how the width of the confidence interval changes when: (a) The sample size increases. (b) The significance level α increases from 0.05 to 0.10. (c) The population standard deviation increases.
- (h) A quality control manager wants to test if the average diameter of manufactured parts differs from the target diameter of 50 mm. The test is conducted at $\alpha = 0.01$ with a sample size of 30 and standard deviation of 2 mm. If the sample mean shifts by 1 mm, discuss how this affects the power of the test. Also, describe how changing α or sample size would affect the confidence interval for the mean diameter.

6.3.5 Questions: Confidence Interval for Difference of Two Means

- (a) Sample A: $n_1 = 100$, $\bar{x}_1 = 65$, $\sigma_1 = 10$
 Sample B: $n_2 = 120$, $\bar{x}_2 = 60$, $\sigma_2 = 12$
 Construct a 95% confidence interval for the difference in means $\mu_1 - \mu_2$.
- (b) A company tests two production lines. Line A ($n = 80$) has an average output time of 25 mins, $\sigma = 2.4$ mins. Line B ($n = 100$) has mean time 26.1 mins, $\sigma = 2.8$ mins. Construct a 90% confidence interval for the difference in mean times.
- (c) Two types of fertilizers are tested on crop yield. Fertilizer A ($n = 60$, $\bar{x} = 2400$ kg/acre, $\sigma = 120$), Fertilizer B ($n = 50$, $\bar{x} = 2250$ kg/acre, $\sigma = 130$). Find a 99% confidence interval for the difference in mean yields.
- (d) A marketing team measures average monthly spending for two customer segments:
 Segment 1: $n = 150$, mean = \$220, $\sigma = 30$
 Segment 2: $n = 140$, mean = \$210, $\sigma = 28$
 Construct a 95% confidence interval for the difference in average spending.

7 Hypothesis Testing: Mean and Difference of Means

7.1 Testing a Single Mean

- Q1.** A sample of 50 students has an average score of 78 with a known population standard deviation of 10. Test whether the population mean is 75 at $\alpha = 0.05$.
- Q2.** A tech recruiter claims that applicants for a new programming bootcamp score lower than 75 on a coding aptitude test. From a random sample of 36 applicants, the average score is 73 with a known population standard deviation of 6. At the 0.05 level of significance, perform the hypothesis test. Compute the p-value and interpret the result.
- Q3.** A manufacturer claims the average lifetime of a battery is 100 hours. A random sample of 32 batteries has a mean life of 96 hours and sample standard deviation of 8 hours. Test the claim at the 0.01 significance level (assume normality).

- Q4.** A company claims its machines produce rods with a mean length of 10 cm. A quality control engineer takes a sample of 30 rods, finds a mean of 9.8 cm, and a sample standard deviation of 0.4 cm. Test the company's claim at $\alpha = 0.05$.
- Q5.** A software company claims that its new algorithm improves processing speed to more than 120 milliseconds per request. A sample of 50 requests is taken, with a mean time of 122 ms and a population standard deviation of 5 ms. At the 0.01 significance level, test the claim using a z-test. Interpret the result.
- Q6.** A factory produces light bulbs with lifetimes normally distributed with mean 1000 hours and standard deviation 100 hours. A quality control test is set up with:

$$H_0 : \mu = 1000, \quad H_a : \mu < 1000,$$

at significance level $\alpha = 0.05$. The test rejects H_0 if the sample mean lifetime $\bar{X}$ from a sample of size 36 is less than 970 hours.

- (a) Calculate the probability of a Type I error.
 - (b) Calculate the probability of a Type II error if the true mean lifetime is 950 hours.
- Q7.** For the same setup in Question 1, find the power of the test when the true mean lifetime is 950 hours.
- Q8.** A hypothesis test is performed to check if a new teaching method improves test scores. The null hypothesis is $H_0 : \mu = 75$ and the alternative is $H_a : \mu > 75$. Suppose the test uses $\alpha = 0.01$, sample size 25, and population standard deviation $\sigma = 10$. If the true mean score is actually 80, calculate the Type II error β and power of the test.
- Q9.** Explain how the width of a confidence interval for a population mean is affected by:
- (a) Increasing the confidence level (e.g., from 95% to 99%)
 - (b) Increasing the sample size
 - (c) Increasing the population standard deviation
- Q10.** A pharmaceutical company tests if a drug decreases cholesterol levels. The test uses $\alpha = 0.05$ and sample size 40 with $\sigma = 15$. The null hypothesis is $H_0 : \mu = 200$, alternative $H_a : \mu < 200$. Calculate the Type II error if the true mean cholesterol level is 185, and discuss how increasing the sample size affects the Type II error.
- Q11.** Discuss the trade-off between Type I and Type II errors. How does changing the significance level α affect these errors and the power of a test?

7.2 Testing the Difference of Two Independent Means

- (a) The mean weight of widgets from machine A is 50g (sample size 40, SD = 5g), and from machine B is 52g (sample size 35, SD = 6g). Test whether there is a significant difference in mean weights at $\alpha = 0.05$.
- (b) Two teaching methods are tested on two independent groups. Group 1 ($n=33$) has a mean score of 85 (SD = 6), and Group 2 ($n=31$) has a mean score of 82 (SD = 5). Test whether the first method leads to higher scores using a one-tailed test at 1% level of significance.
- (c) Two data centers report different average response times. A sample of 100 requests from Data Center A yields a mean response time of 210 ms with a population standard deviation of 15 ms. A sample of 120 requests from Data Center B yields a mean of 215 ms with a population standard deviation of 20 ms. Test at the 0.05 significance level whether there is a significant difference in mean response times.

7.3 Hypothesis Testing: Regression Parameters

7.3.1 Simple Linear Regression (Large Sample)

- Q1.** In the simple linear regression model $y = \beta_0 + \beta_1 x + \varepsilon$, define the null and alternative hypotheses for testing the slope β_1 using a z-test.
- Q2.** A simple regression analysis yields:

$$\hat{y} = 1.2 + 3.4x, \quad SE_{\beta_1} = 0.5$$

Assume $n = 1000$ (a large sample). Test whether β_1 is significantly different from zero at the 5% significance level using the z-test.

7.3.2 Multiple Linear Regression (Large Sample)

- (a) For the multiple regression model:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \varepsilon$$

state the null and alternative hypotheses to test if $\beta_2 = 0$ using a z-test.

- (b) A multiple regression model output shows:

$$\hat{\beta}_2 = 0.75, \quad SE_{\beta_2} = 0.25, \quad n = 1200$$

Use a z-test to determine whether x_2 significantly affects y at the 1% significance level.

- (c) A digital marketing team develops a multiple linear regression model to predict the weekly number of purchases (y) from an e-commerce website using the following inputs:
- x_1 : Weekly advertising budget on social media (in \$1000s)
 - x_2 : Website traffic (in 1000s of visits)
 - x_3 : Average session duration (in minutes)

The estimated regression equation is:

$$\hat{y} = 10 + 1.8x_1 + 2.5x_2 + 0.1x_3$$

The standard errors of the estimated coefficients based on a sample size of $n = 1500$ are:

$$SE_{\beta_1} = 0.4, \quad SE_{\beta_2} = 0.6, \quad SE_{\beta_3} = 0.3$$

Test whether each predictor significantly affects the weekly number of purchases at the 1% significance level using a z-test.

8 Logistic Regression

8.1 Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

8.1.1 Symmetry:

Show that $\sigma(-z) = 1 - \sigma(z)$

$$\begin{aligned}\sigma(-z) &= \frac{1}{1 + e^z}, \quad \sigma(z) = \frac{1}{1 + e^{-z}} \\ 1 - \sigma(z) &= 1 - \frac{1}{1 + e^{-z}} = \frac{e^{-z}}{1 + e^{-z}} = \frac{1}{1 + e^z} = \sigma(-z)\end{aligned}$$

8.1.2 Inverse:

Solve $p = \sigma(z)$ for z

$$\begin{aligned}p &= \frac{1}{1 + e^{-z}} \Rightarrow 1 + e^{-z} = \frac{1}{p} \Rightarrow e^{-z} = \frac{1-p}{p} \\ \Rightarrow -z &= \log\left(\frac{1-p}{p}\right) \Rightarrow z = \log\left(\frac{p}{1-p}\right)\end{aligned}$$

8.1.3 Derivative:

$$\begin{aligned}\sigma(z) &= \frac{1}{1 + e^{-z}} \Rightarrow \frac{d\sigma}{dz} = \frac{e^{-z}}{(1 + e^{-z})^2} \\ &= \left(\frac{1}{1 + e^{-z}}\right) \left(1 - \frac{1}{1 + e^{-z}}\right) = \sigma(z)(1 - \sigma(z))\end{aligned}$$

8.2 Introduction to Logistic Regression

Purpose: Logistic regression is used to model the probability of a binary outcome (e.g., success/failure, yes/no, 0/1) based on one or more predictor variables.

Model Form:

$$P(Y = 1 \mid X) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \dots + \beta_j X_j)}}$$

Log-Odds (Logit) Form:

$$\log\left(\frac{P(Y = 1)}{1 - P(Y = 1)}\right) = \beta_0 + \beta_1 X_1 + \dots + \beta_j X_j$$

Why Use It?

- Predicts probability while keeping values in $[0, 1]$
- Interpretable via odds and log-odds
- Suitable for classification problems

8.2.1 Why Logistic Regression Has This Form

Problem: We want to model a probability $P(Y = 1 | X = x)$, but a linear model like $\beta_0 + \beta_1 x$ can take values outside the interval $[0, 1]$.

Solution: Use the *logit function*, which maps probabilities to the entire real line:

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

Model:

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x \quad \Rightarrow \quad p = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

Why this form?

- Ensures $0 < p < 1$ for all x
- Logit is the inverse of the sigmoid function
- Interpretable: coefficients show how log-odds change with x

Meaning of Logistic Coefficient β_j **Model:**

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1 + \dots + \beta_j X_j$$

Interpretation:

- Each β_j represents the change in **log-odds** of $Y = 1$ for a one-unit increase in X_j .
- Exponentiating β_j gives the **odds ratio**:

$$\text{Odds Ratio} = e^{\beta_j}$$

- If $\beta_j > 0$, the odds of $Y = 1$ increase with X_j .
- If $\beta_j < 0$, the odds decrease.

8.2.2 Decision Boundary in Logistic Regression

$$P(Y = 1 | \mathbf{X}) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \dots + \beta_j X_j)}}$$

Decision Boundary: Predicted class changes at the threshold where

$$P(Y = 1 | \mathbf{X}) = 0.5$$

Start with:

$$\frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \dots + \beta_j X_j)}} = 0.5$$

Solve:

$$\begin{aligned} 1 + e^{-(\beta_0 + \beta_1 X_1 + \dots + \beta_j X_j)} &= 2 \\ e^{-(\beta_0 + \dots)} &= 1 \quad \Rightarrow \quad \beta_0 + \beta_1 X_1 + \dots + \beta_j X_j = 0 \end{aligned}$$

Hence Decision Boundary:

$$\beta_0 + \beta_1 X_1 + \dots + \beta_j X_j = 0$$

- It is a **hyperplane** in the input space.

- Separates predictions where $P(Y = 1) > 0.5$ from those where $P(Y = 1) < 0.5$.
- Equivalent to a linear classifier boundary.

Questions

- (a) A logistic regression model outputs a probability of 0.87 that a given email is spam. The threshold to classify an email as spam is 0.5. Should the email be classified as spam or not spam? Justify your answer.
- (b) A logistic regression model outputs that the odds of a user clicking on an ad are 7:3. What is the corresponding probability?
- (c) A logistic regression model predicts whether a user will click on a recommended item ($y = 1$), based on the item popularity score (x_1) and user activity level (x_2). The model is:

$$P(y = 1 \mid x_1, x_2) = \frac{1}{1 + e^{-(-2 + 0.5x_1 + 0.1x_2)}}$$

Calculate the probability that a user with popularity score 6 and activity level 40 clicks on the recommendation.

- (d) A logistic regression model is used to predict whether a user will buy a product ($y = 1$) or not ($y = 0$), based on the time they spent on the website (in minutes), denoted by x . The model is:

$$P(y = 1 \mid x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

Suppose the parameters of the model are:

$$\beta_0 = -2, \quad \beta_1 = 0.5$$

You are given the following training data:

x	1	2	3	4
y	0	0	1	1

Compute the log-likelihood function for this dataset under the given model.

- (e) Derive the gradient (first derivative) of the log-likelihood function with respect to β_0 and β_1 .
- (f) A company is developing a spam detection system using logistic regression. The model predicts whether an email is spam ($y = 1$) or not spam ($y = 0$) based on a single feature x , which represents the number of spam-related keywords found in the email.

The logistic regression model is given by:

$$P(y = 1 \mid x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

with parameters:

$$\beta_0 = 0.2, \quad \beta_1 = 1.5$$

The following emails are analyzed, with the number of spam keywords found in each:

$$x = \{0, 1, 2\}$$

- (a) Calculate the predicted probability $P(y = 1 \mid x)$ for each value of x .
- (b) Classify each email as spam or not spam using a threshold of 0.7. If $P(y = 1) \geq 0.7$, classify as spam ($y = 1$); otherwise, classify as not spam ($y = 0$).

- (g) A logistic regression model is used to predict whether a customer will purchase a product ($y = 1$) based on their browsing time x (in minutes). The model is given by:

$$P(y = 1 \mid x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

Browsing Time (x)	1.5	2.5	3.5	4.5
Purchase (y)	0	1	1	0

- (a) Write the likelihood function $L(\beta_0, \beta_1)$ based on the given data.
- (b) Compute the likelihood value when $\beta_0 = 0$ and $\beta_1 = 1$.